

# EE559 Project Assignment

April 28, 2022

## 1 *Student Performance Dataset / Classification*

### *Imports*

```
[ ]: from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import NearestCentroid
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import random
import seaborn as sns
```

```
[ ]: import warnings
warnings.filterwarnings('ignore')
```

### *Functions*

```
[ ]: # Fucntion to convert scores to grades
def get_grade(score):
    if score <= 9:
        grade = 0 #F
    elif 10<=score<=11:
        grade = 1 #D
    elif 12<=score<=13:
        grade = 2 #C
    elif 14<=score<=15:
        grade = 3 #B
    elif score>=16:
        grade = 4 #A
```

```

    return grade

# Trivial classifier - randomly outputs class labels with probability based on
↳class priors
def trivial_classifier(Y_train, data):
    y_grades_train = list()
    for i in range(len(Y_train)):
        y_grades_train.append(get_grade(Y_train[i]))
    y_grades_train = np.array(y_grades_train)
    weights = [np.count_nonzero(y_grades_train == 0)/len(Y_train)*100, np.
↳count_nonzero(y_grades_train == 1)/len(Y_train)*100,
    np.count_nonzero(y_grades_train == 2)/len(Y_train)*100, np.
↳count_nonzero(y_grades_train == 3)/len(Y_train)*100,
    np.count_nonzero(y_grades_train == 4)/len(Y_train)*100]
    y_pred = list()
    for i in range(len(data)):
        pred_class = 0
        for j in range(10):
            gradelist = [0, 1, 2, 3, 4]
            pred_class += random.choices(gradelist, weights=weights)[0]
        y_pred.append(round(pred_class/10))
    return np.array(y_pred)

# Baseline Model - Nearest Means Classifier
def nearestMeansClassifier(X_train, Y_train, X_test):
    y_grades_train = list()
    for i in range(len(Y_train)):
        y_grades_train.append(get_grade(Y_train[i]))
    y_grades_train = np.array(y_grades_train)
    clf = NearestCentroid()
    clf.fit(X_train, y_grades_train)
    Y_test_grades_pred = np.zeros(len(X_test))
    for i in range(len(X_test)):
        Y_test_grades_pred[i] = clf.predict(X_test[i].reshape(1, -1))

    return Y_test_grades_pred

# Performance Measures
def get_performance(y, y_pred):
    macro_f1_score = f1_score(y, y_pred, average = 'macro')
    print("The macro F1 score for the classifier is - ", macro_f1_score)
    accuracy = f1_score(y, y_pred, average = 'micro')
    print("The accuracy for the classifier is - ", accuracy)
    cf_matrix = confusion_matrix(y, y_pred)
    labels = ['F', 'D', 'C', 'B', 'A']
    ax = sns.heatmap(cf_matrix, annot=True)
    ax.set_xticklabels(labels)

```

```
ax.set_yticklabels(labels)
plt.show()
```

```
[ ]: def KNNClassifier(X, y, X_test, Y_test_actual):
    params = {
        'n_neighbors' : [2,3,4,5],
        'algorithm' : ['ball_tree', 'kd_tree', 'brute', 'brute_force'],
        'auto'],
        'leaf_size' : [10, 20, 30, 40, 50],
        'weights' : ['uniform', 'distance']
    }
    neigh = KNeighborsClassifier()
    clf = GridSearchCV(estimator=neigh, param_grid=params, cv=5)
    clf.fit(X, y)
    print('Best parameters from Cross Validation: ', clf.best_params_)
    #print('\n')
    y_train_pred = clf.predict(X)
    accuracy = f1_score(y, y_train_pred, average = 'micro')
    print('Cross Validation Best Score', accuracy)
    print('\n')
    Y_test_grades_pred = clf.predict(X_test)
    get_performance(Y_test_actual, Y_test_grades_pred)

def LogitClassifier(X, y, X_test, Y_test_actual):
    params= {
        'tol' : [1e-5, 0.5*1e-4, 1e-4, 2*1e-4, 1e-3],
        'solver': ['newton-cg', 'lbfgs', 'liblinear']
    }
    logreg = LogisticRegression(max_iter=100000)
    clf = GridSearchCV(estimator=logreg, param_grid=params, cv=5)
    clf.fit(X, y)
    print('Best parameters from Cross Validation: ', clf.best_params_)
    y_train_pred = clf.predict(X)
    accuracy = f1_score(y, y_train_pred, average = 'micro')
    print('Cross Validation Best Score', accuracy)
    Y_test_grades_pred = clf.predict(X_test)
    get_performance(Y_test_actual, Y_test_grades_pred)

def KernelSVMClassifier(X, y, X_test, Y_test_actual):
    params= {
        'C' : [0.8, 0.9, 1.0, 1.1, 1.2],
        'kernel' : ['linear', 'poly', 'rbf', 'sigmoid'],
        'tol' : [1e-5, 1e-4, 1e-3],
        'gamma' : ['scale', 'auto'],
        'class_weight' : ['balanced', None]
    }
```

```

svc = SVC()
clf = GridSearchCV(estimator=svc, param_grid=params, cv=5)
clf.fit(X, y)
print('Best parameters from Cross Validation: ', clf.best_params_)
y_train_pred = clf.predict(X)
accuracy = f1_score(y, y_train_pred, average = 'micro')
print('Cross Validation Best Score', accuracy)
Y_test_grades_pred = clf.predict(X_test)
get_performance(Y_test_actual, Y_test_grades_pred)

def MLPPerceptronClassifier(X, y, X_test, Y_test_actual):
    params = {
        'activation' : ['tanh', 'relu', 'identity', 'logistic'],
        'solver' : ['sgd', 'adam', 'lbfgs'],
        'alpha' : [0.0001, 0.05],
        'learning_rate' : ['constant', 'adaptive'],
    }
    mlp = MLPClassifier(random_state=1, max_iter=2000)
    clf = GridSearchCV(estimator=mlp, param_grid=params, cv=5)
    clf.fit(X, y)
    print('Best parameters from Cross Validation: ', clf.best_params_)
    print('Cross Validation Best Score', clf.best_score_)
    Y_test_grades_pred = clf.predict(X_test)
    get_performance(Y_test_actual, Y_test_grades_pred)
    pass

```

Read in the data

```
[ ]: train_df = pd.read_csv('data/student_performance_train.csv')
#dataset_train = train_df.to_numpy()
```

```
[ ]: train_df.columns
```

```
[ ]: Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
'Walc', 'health', 'absences', 'G1', 'G2', 'G3'],
dtype='object')
```

## 1.1 Problem 1

- Predict first-period academic performance without any prior academic performance data: remove the G2 and G3 columns from the original dataset, then predict G1.

Removing categorical non-binary features and grades.

```
[ ]: selected_columns = train_df.loc[:, ~train_df.columns.isin(['Mjob', 'Fjob', 'reason', 'gaurdian', 'G1', 'G2', 'G3'])]
binary_vals = pd.get_dummies(selected_columns)
X_train = binary_vals.to_numpy() #Converting to numpy array for easier processing
```

```
[ ]: y_train = train_df.loc[:, train_df.columns.isin(['G1'])]
Y_train = y_train['G1'].to_numpy()
y_grades_train = list()
for i in range(len(Y_train)):
    y_grades_train.append(get_grade(Y_train[i]))
y_grades_train = np.array(y_grades_train)
```

Reading in test data and processing it

```
[ ]: test_df = pd.read_csv('data/student_performance_test.csv')
selected_columns = test_df.loc[:, ~test_df.columns.isin(['Mjob', 'Fjob', 'reason', 'gaurdian', 'G1', 'G2', 'G3'])]
binary_vals = pd.get_dummies(selected_columns)
X_test = binary_vals.to_numpy()
```


```
[ ]: y_test = test_df.loc[:, test_df.columns.isin(['G1'])] #Rerun this cell to get back original values of Y_test
Y_test = y_test['G1'].to_numpy()
y_grades_test_actual = list()
for i in range(len(Y_test)):
    y_grades_test_actual.append(get_grade(Y_test[i]))
Y_test_grades_actual = np.array(y_grades_test_actual)
```

### 1.1.1 Trivial System

```
[ ]: Y_test_grades_pred = trivial_classifier(Y_train, X_test)
get_performance(Y_test_grades_actual, Y_test_grades_pred)
```

The macro F1 score for the classifier is - 0.1511294694720537

The accuracy for the classifier is - 0.2883435582822086




### 1.1.2 Reference System - Nearest Means

```
[ ]: Y_test_grades_pred = nearestMeansClassifier(X_train, Y_train, X_test)
get_performance(Y_test_grades_actual, Y_test_grades_pred)
```

The macro F1 score for the classifier is - 0.2339098179522256

The accuracy for the classifier is - 0.245398773006135



Normalizing Data and encoding categorical data

```
[ ]: selected_columns = train_df.loc[:, ~train_df.columns.isin(['G1', 'G2', 'G3'])]
      binary_vals = pd.get_dummies(selected_columns)
      X_train = binary_vals.to_numpy()
      pipe = Pipeline([('scale', StandardScaler())])
      X_train_scaled = pipe.fit_transform(X_train)
```


```
[ ]: test_df = pd.read_csv('data/student_performance_test.csv')
      selected_columns = test_df.loc[:, ~test_df.columns.isin(['G1', 'G2', 'G3'])]
      binary_vals = pd.get_dummies(selected_columns)
      X_test = binary_vals.to_numpy()
      pipe = Pipeline([('scale', StandardScaler())])
      X_test_scaled = pipe.fit_transform(X_test)
```

### 1.1.3 Approach 1: K Nearest Neighbors

```
[ ]: KNNClassifier(X_train_scaled, y_grades_train, X_test_scaled, ↵
      ↪y_grades_test_actual)
```

Best parameters from Cross Validation: {'algorithm': 'ball\_tree', 'leaf\_size': 10, 'n\_neighbors': 5, 'weights': 'uniform'}  
Cross Validation Best Score 0.5390946502057613

The macro F1 score for the classifier is - 0.2555260509888779  
The accuracy for the classifier is - 0.3312883435582822



#### 1.1.4 Approach 2 - Logistic Regression

```
[ ]: LogitClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪y_grades_test_actual)
```


Best parameters from Cross Validation: {'solver': 'liblinear', 'tol': 1e-05}

Cross Validation Best Score 0.5

The macro F1 score for the classifier is - 0.2662989050348655

The accuracy for the classifier is - 0.31901840490797545





### 1.1.5 Approach 3 - Kernel Support Vector Machine


```
[ ]: KernelSVMClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪ y_grades_test_actual)
```

Best parameters from Cross Validation: {'C': 1.2, 'class\_weight': None, 'gamma': 'scale', 'kernel': 'rbf', 'tol': 1e-05}

Cross Validation Best Score 0.8086419753086419

The macro F1 score for the classifier is - 0.29536184623102973

The accuracy for the classifier is - 0.3803680981595092



### 1.1.6 Approach 4 - Multi-Layer Perceptron

```
[ ]: MLPPerceptronClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪ y_grades_test_actual)
```

Best parameters from Cross Validation: {'activation': 'logistic', 'alpha': 0.05, 'learning\_rate': 'constant', 'solver': 'sgd'}

Cross Validation Best Score 0.3560067325899432

The macro F1 score for the classifier is - 0.27350277460658046

The accuracy for the classifier is - 0.34355828220858897



## 1.2 Problem 2

- Predict final-period academic performance without any prior academic performance data: remove the G1 and G2 columns from the original dataset, then predict G3.

Removing categorical non-binary features and grades.

```
[ ]: selected_columns = train_df.loc[:, ~train_df.columns.isin(['Mjob', 'Fjob', 'reason', 'gaurdian', 'G1', 'G2', 'G3'])]
      binary_vals = pd.get_dummies(selected_columns)
      X_train = binary_vals.to_numpy() #Converting to numpy array for easier processing
      y_train = train_df.loc[:, train_df.columns.isin(['G3'])] #Getting corresponding labels
      Y_train = y_train['G3'].to_numpy()
      y_grades_train = list()
      for i in range(len(Y_train)):
          y_grades_train.append(get_grade(Y_train[i]))
      y_grades_train = np.array(y_grades_train)
```

Reading in test data and processing it

```
[ ]: test_df = pd.read_csv('data/student_performance_test.csv')
      selected_columns = test_df.loc[:, ~test_df.columns.isin(['Mjob', 'Fjob', 'reason', 'gaurdian', 'G1', 'G2', 'G3'])]
```

```

binary_vals = pd.get_dummies(selected_columns)
X_test = binary_vals.to_numpy() #Converting to numpy array for easier processing

y_test = test_df.loc[:, test_df.columns.isin(['G3'])] #Getting corresponding
↪ labels
Y_test = y_test['G3'].to_numpy()
y_grades_test_actual = list()
for i in range(len(Y_test)):
    y_grades_test_actual.append(get_grade(Y_test[i]))
Y_test_grades_actual = np.array(y_grades_test_actual)

```

### 1.2.1 Trivial System


```

[ ]: Y_test_grades_pred = trivial_classifier(Y_train, X_test)
get_performance(Y_test_grades_actual, Y_test_grades_pred)

```

The macro F1 score for the classifier is - 0.16115332513972497

The accuracy for the classifier is - 0.2883435582822086



### 1.2.2 Reference System - Nearest Means

```

[ ]: Y_test_grades_pred = nearestMeansClassifier(X_train, Y_train, X_test)
get_performance(Y_test_grades_actual, Y_test_grades_pred)

```

The macro F1 score for the classifier is - 0.25040487766154157

The accuracy for the classifier is - 0.26380368098159507




### 1.2.3 Approach 1: K Nearest Neighbors

```
[ ]: KNNClassifier(X_train_scaled, y_grades_train, X_test_scaled, y_grades_test_actual)
```

Best parameters from Cross Validation: {'algorithm': 'ball\_tree', 'leaf\_size': 10, 'n\_neighbors': 3, 'weights': 'uniform'}  
Cross Validation Best Score 0.6049382716049383

The macro F1 score for the classifier is - 0.3245536127709303

The accuracy for the classifier is - 0.3619631901840491



#### 1.2.4 Approach 2 - Logistic Regression


```
[ ]: LogitClassifier(X_train_scaled, y_grades_train, X_test_scaled,
                    ↪y_grades_test_actual)
```

Best parameters from Cross Validation: {'solver': 'liblinear', 'tol': 1e-05}

Cross Validation Best Score 0.5102880658436214

The macro F1 score for the classifier is - 0.3159904678197361

The accuracy for the classifier is - 0.3374233128834356



### 1.2.5 Approach 3 - Kernel Support Vector Machine


```
[ ]: KernelSVMClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪ y_grades_test_actual)
```

Best parameters from Cross Validation: {'C': 1.0, 'class\_weight': 'balanced', 'gamma': 'scale', 'kernel': 'rbf', 'tol': 1e-05}

Cross Validation Best Score 0.8065843621399177

The macro F1 score for the classifier is - 0.3444695272837712

The accuracy for the classifier is - 0.3496932515337423



### 1.2.6 Approach 4 - Multi Layer Perceptron

```
[ ]: MLPPerceptromClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪y_grades_test_actual)
```

Best parameters from Cross Validation: {'activation': 'relu', 'alpha': 0.0001, 'learning\_rate': 'constant', 'solver': 'adam'}

Cross Validation Best Score 0.3496949295181991

The macro F1 score for the classifier is - 0.339720311316846

The accuracy for the classifier is - 0.34355828220858897





### 1.3 Problem 3

- Predict final academic performance using all available prior academic performance data: Keep G1 and G2 columns inside the dataset as features, then predict G3.

Removing categorical non-binary features.

```
[ ]: selected_columns = train_df.loc[:, ~train_df.columns.isin(['Mjob', 'Fjob', 'reason', 'gaurdian', 'G3'])]
      binary_vals = pd.get_dummies(selected_columns)
      X_train = binary_vals.to_numpy() #Converting to numpy array for easier processing
      y_train = train_df.loc[:, train_df.columns.isin(['G3'])] #Getting corresponding labels
      Y_train = y_train['G3'].to_numpy()
      y_grades_train = list()
      for i in range(len(Y_train)):
          y_grades_train.append(get_grade(Y_train[i]))
      y_grades_train = np.array(y_grades_train)
```

Reading in test data and processing it

```
[ ]: test_df = pd.read_csv('data/student_performance_test.csv')
      selected_columns = test_df.loc[:, ~test_df.columns.isin(['Mjob', 'Fjob', 'reason', 'gaurdian', 'G3'])]
```

```

binary_vals = pd.get_dummies(selected_columns)
X_test = binary_vals.to_numpy() #Converting to numpy array for easier processing

y_test = test_df.loc[:, test_df.columns.isin(['G3'])] #Getting corresponding
↪ labels
Y_test = y_test['G3'].to_numpy()
y_grades_test_actual = list()
for i in range(len(Y_test)):
    y_grades_test_actual.append(get_grade(Y_test[i]))
Y_test_grades_actual = np.array(y_grades_test_actual)

```

### 1.3.1 Trivial System


```

[ ]: Y_test_grades_pred = trivial_classifier(Y_train, X_test)
get_performance(Y_test_grades_actual, Y_test_grades_pred)

```

The macro F1 score for the classifier is - 0.1182763337893297

The accuracy for the classifier is - 0.2147239263803681



### 1.3.2 Reference System - Nearest Means


```

[ ]: Y_test_grades_pred = nearestMeansClassifier(X_train, Y_train, X_test)
get_performance(Y_test_grades_actual, Y_test_grades_pred)

```

The macro F1 score for the classifier is - 0.6150584274113686

The accuracy for the classifier is - 0.6012269938650306



Normalizing data and encoding categorical data

```
[ ]: selected_columns = train_df.loc[:, ~train_df.columns.isin(['G3'])]
      binary_vals = pd.get_dummies(selected_columns)
      X_train = binary_vals.to_numpy()
      pipe = Pipeline([('scale', StandardScaler())])
      X_train_scaled = pipe.fit_transform(X_train)
```

```
[ ]: test_df = pd.read_csv('data/student_performance_test.csv')
      selected_columns = test_df.loc[:, ~test_df.columns.isin(['G3'])]
      binary_vals = pd.get_dummies(selected_columns)
      X_test = binary_vals.to_numpy()
      pipe = Pipeline([('scale', StandardScaler())])
      X_test_scaled = pipe.fit_transform(X_test)
```

### 1.3.3 Approach 1: K Nearest Neighbors


```
[ ]: KNNClassifier(X_train_scaled, y_grades_train, X_test_scaled,
                  ↪y_grades_test_actual)
```

Best parameters from Cross Validation: {'algorithm': 'ball\_tree', 'leaf\_size': 10, 'n\_neighbors': 5, 'weights': 'uniform'}

Cross Validation Best Score 0.6069958847736625

The macro F1 score for the classifier is - 0.3943313305498517

The accuracy for the classifier is - 0.4049079754601227



### 1.3.4 Approach 2 - Logistic Regression


```
[ ]: LogitClassifier(X_train_scaled, y_grades_train, X_test_scaled, y_grades_test_actual)
```

Best parameters from Cross Validation: {'solver': 'newton-cg', 'tol': 1e-05}

Cross Validation Best Score 0.8045267489711934

The macro F1 score for the classifier is - 0.7222694026818367

The accuracy for the classifier is - 0.7361963190184049



### 1.3.5 Approach 3 - Kernel Support Vector Machine


```
[ ]: KernelSVMClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪ y_grades_test_actual)
```

Best parameters from Cross Validation: {'C': 1.2, 'class\_weight': None, 'gamma': 'scale', 'kernel': 'linear', 'tol': 1e-05}

Cross Validation Best Score 0.8559670781893005

The macro F1 score for the classifier is - 0.707249254649567

The accuracy for the classifier is - 0.7239263803680982



### 1.3.6 Approach 4 - Multi Layer Perceptron

```
[ ]: MLPPerceptronClassifier(X_train_scaled, y_grades_train, X_test_scaled,
    ↪ y_grades_test_actual)
```

Best parameters from Cross Validation: {'activation': 'relu', 'alpha': 0.0001, 'learning\_rate': 'constant', 'solver': 'lbfgs'}

Cross Validation Best Score 0.615190406059331

The macro F1 score for the classifier is - 0.6325642922845869

The accuracy for the classifier is - 0.6257668711656442

